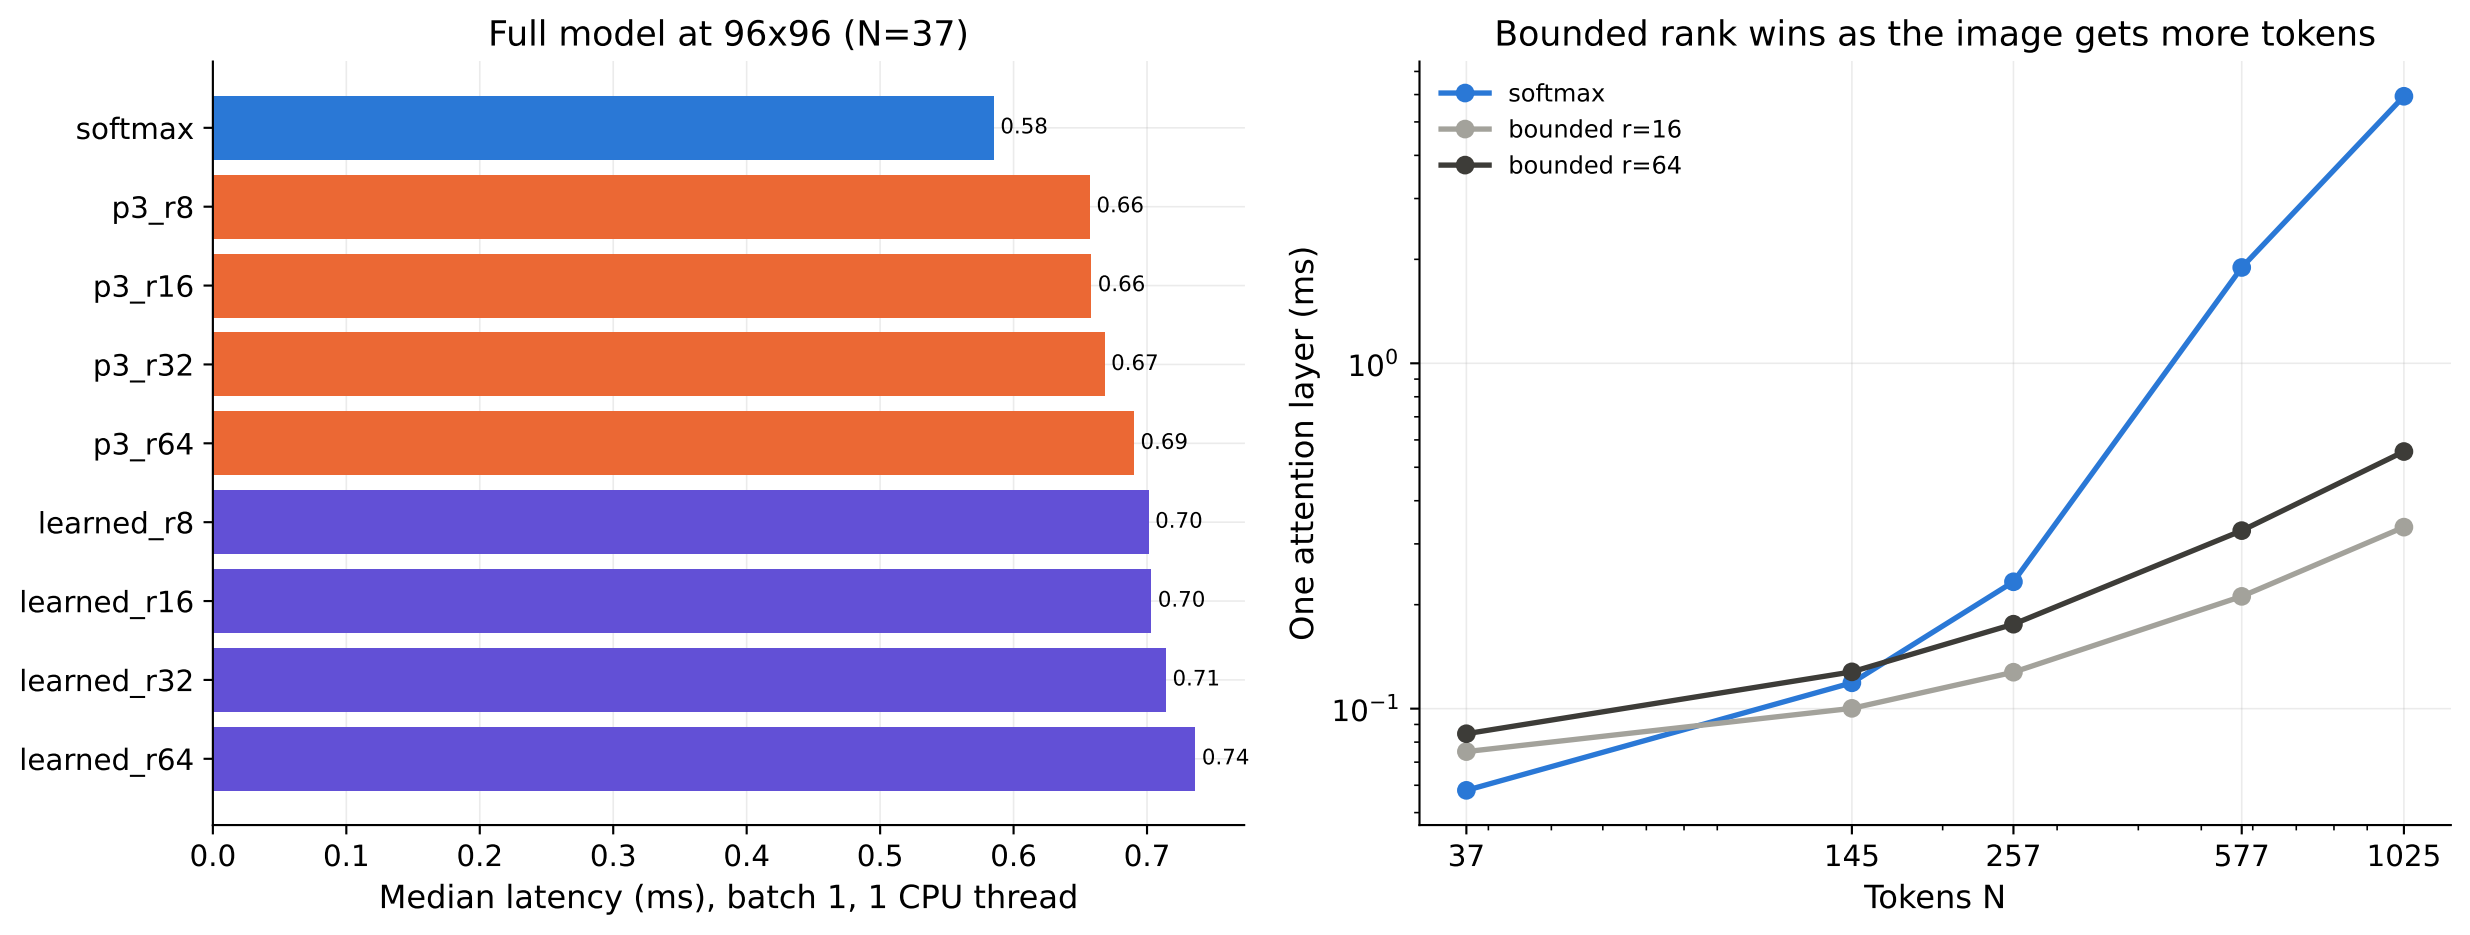
Full model at 96x96 (N=37)
softmax
p3_r8
p3_r16
p3_r32
p3_r64
learned_r8
learned_r16
learned_r32
learned_r64
0.58
0.66
0.66
0.67
0.69
0.70
0.70
0.71
0.74
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
Median latency (ms), batch 1, 1 CPU thread
Bounded rank wins as the image gets more tokens
softmax
bounded r=16
bounded r=64
One attention layer (ms)
10^{0}
10^{-1}
37
145
257
577
1025
Tokens N