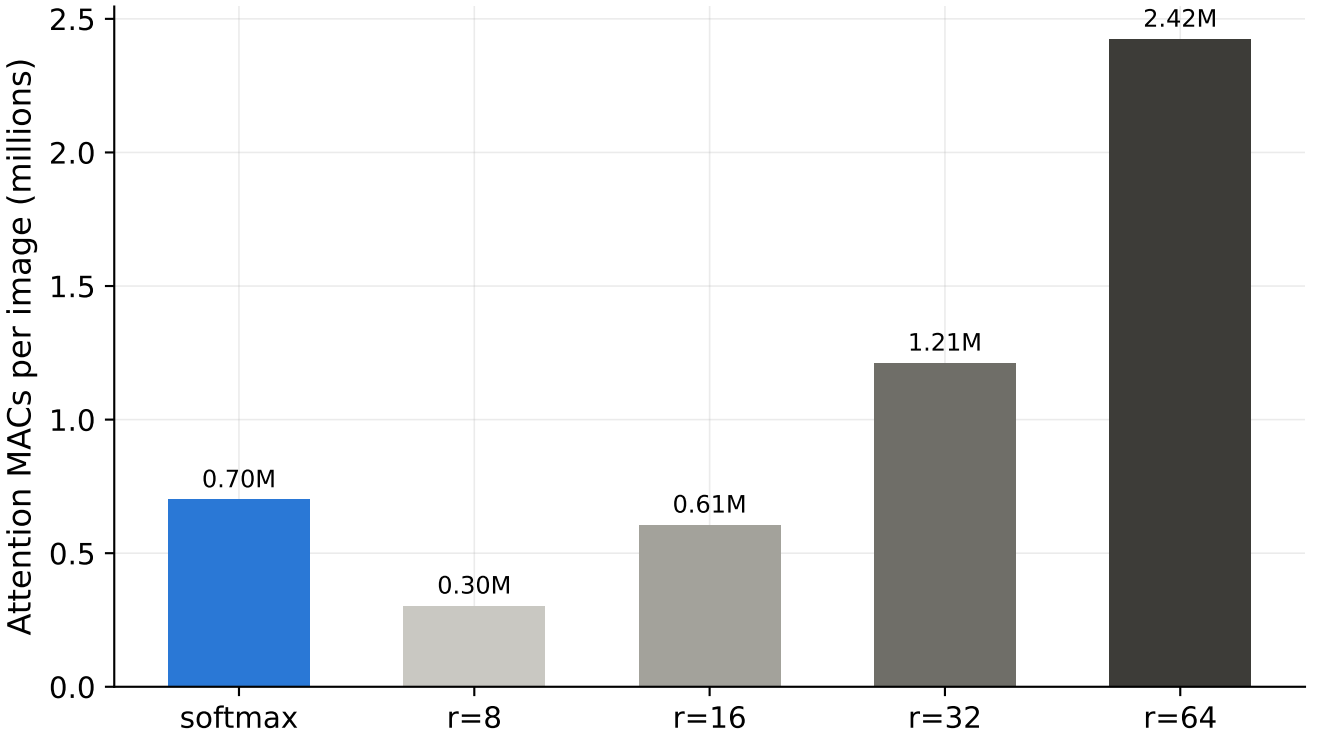
Attention cost at N=37 tokens: bounded rank is cheaper only for r < N/2
Attention MACs per image (millions)
2.5
2.0
1.5
1.0
0.5
0.0
0.70M
0.30M
0.61M
1.21M
2.42M
softmax
r=8
r=16
r=32
r=64